

hypothesis-testing-using-python

July 10, 2024

Hypothesis Testing is a statistical method used to make inferences or decisions about a population based on sample data. It starts with a null hypothesis (H_0), which represents a default stance or no effect, and an alternative hypothesis (H_1 or H_a), which represents what we aim to prove or expect to find. The process involves using sample data to determine whether to reject the null hypothesis in favor of the alternative hypothesis, based on the likelihood of observing the sample data under the null hypothesis. Hypothesis Testing: Process We Can Follow

So, Hypothesis Testing is a fundamental process in data science for making data-driven decisions and inferences about populations based on sample data. Below is the process we can follow for the task of Hypothesis Testing:

1. Gather the necessary data required for the hypothesis test.
2. Define Null (H_0) and Alternative Hypothesis (H_1 or H_a).
3. Choose the Significance Level (α), which is the probability of rejecting the null hypothesis when it is true.
4. Select the appropriate statistical tests. Examples include t-tests for comparing means, chi-square tests for categorical data, and ANOVA for comparing means across more than two groups.
5. Perform the chosen statistical test on your data.
6. Determine the p-value and interpret the results of your statistical tests

```
[4]: import pandas as pd
from scipy.stats import ttest_ind

df = pd.read_csv("website_ab_test.csv")
df.head()
```

```
[4]:
```

	Theme	Click Through Rate	Conversion Rate	Bounce Rate	\
0	Light Theme	0.054920	0.282367	0.405085	
1	Light Theme	0.113932	0.032973	0.732759	
2	Dark Theme	0.323352	0.178763	0.296543	
3	Light Theme	0.485836	0.325225	0.245001	
4	Light Theme	0.034783	0.196766	0.765100	

	Scroll_Depth	Age	Location	Session_Duration	Purchases	Added_to_Cart
0	72.489458	25	Chennai	1535	No	Yes
1	61.858568	19	Pune	303	No	Yes
2	45.737376	47	Chennai	563	Yes	Yes
3	76.305298	58	Pune	385	Yes	No

4	48.927407	25	New Delhi	1437	No	No
---	-----------	----	-----------	------	----	----

So, the dataset is based on the performance of two themes on a website. Our task is to find which theme performs better using Hypothesis Testing. Let's go through the summary of the dataset, including the number of records, the presence of missing values, and basic statistics for the numerical columns:

```
[8]: # dataset summary
summary = {
    'Number of Records': df.shape[0],
    'Number of Columns': df.shape[1],
    'Missing Values': df.isnull().sum(),
    'Numerical Columns Summary':
        df.describe()
}

summary
```

```
[8]: {'Number of Records': 1000,
      'Number of Columns': 10,
      'Missing Values': Theme          0
      Click Through Rate    0
      Conversion Rate       0
      Bounce Rate           0
      Scroll_Depth          0
      Age                   0
      Location              0
      Session_Duration      0
      Purchases             0
      Added_to_Cart         0
      dtype: int64,
      'Numerical Columns Summary':
          Click Through Rate  Conversion Rate  Bounce
Rate Scroll_Depth \
count      1000.000000      1000.000000      1000.000000      1000.000000
mean         0.256048         0.253312         0.505758         50.319494
std          0.139265         0.139092         0.172195         16.895269
min          0.010767         0.010881         0.200720         20.011738
25%          0.140794         0.131564         0.353609         35.655167
50%          0.253715         0.252823         0.514049         51.130712
75%          0.370674         0.373040         0.648557         64.666258
max          0.499989         0.498916         0.799658         79.997108

          Age  Session_Duration
count  1000.000000      1000.000000
mean    41.528000       924.999000
std     14.114334       508.231723
min     18.000000       38.000000
```

25%	29.000000	466.500000
50%	42.000000	931.000000
75%	54.000000	1375.250000
max	65.000000	1797.000000 }

The dataset contains 1,000 records across 10 columns, with no missing values. Here's a quick summary of the numerical columns:

1. Click Through Rate: Ranges from about 0.01 to 0.50 with a mean of approximately 0.26.
2. Conversion Rate: Also ranges from about 0.01 to 0.50 with a mean close to the Click Through Rate, approximately 0.25.
3. Bounce Rate: Varies between 0.20 and 0.80, with a mean around 0.51.
4. Scroll Depth: Shows a spread from 20.01 to nearly 80, with a mean of 50.32.
5. Age: The age of users ranges from 18 to 65 years, with a mean age of about 41.5 years.
6. Session Duration: This varies widely from 38 seconds to nearly 1800 seconds (30 minutes), with a mean session duration of approximately 925 seconds (about 15 minutes).

Now, let's move on to comparing the performance of both themes based on the provided metrics. We'll look into the average Click Through Rate, Conversion Rate, Bounce Rate, and other relevant metrics for each theme. Afterwards, we can perform hypothesis testing to identify if there's a statistically significant difference between the themes:

```
[15]: numeric_cols = df.select_dtypes(include='number').columns
      # grouping data by theme and calculating mean values for the metrics
      theme_performance = df.groupby('Theme')[numeric_cols].mean()

      # sorting the data by conversion rate for a better comparison
      theme_performance_sorted = theme_performance.sort_values(by='Conversion Rate',
      ↪ascending=False)

      print(theme_performance_sorted)
```

	Click Through Rate	Conversion Rate	Bounce Rate	Scroll_Depth \
Theme				
Light Theme	0.247109	0.255459	0.499035	50.735232
Dark Theme	0.264501	0.251282	0.512115	49.926404

	Age	Session_Duration
Theme		
Light Theme	41.734568	930.833333
Dark Theme	41.332685	919.482490

The comparison between the Light Theme and Dark Theme on average performance metrics reveals the following insights:

1. Click Through Rate (CTR): The Dark Theme has a slightly higher average CTR (0.2645) compared to the Light Theme (0.2471).
2. Conversion Rate: The Light Theme leads with a marginally higher average Conversion Rate (0.2555) compared to the Dark Theme (0.2513).

3. Bounce Rate: The Bounce Rate is slightly higher for the Dark Theme (0.5121) than for the Light Theme (0.4990).
4. Scroll Depth: Users on the Light Theme scroll slightly further on average (50.74%) compared to those on the Dark Theme (49.93%).
5. Age: The average age of users is similar across themes, with the Light Theme at approximately 41.73 years and the Dark Theme at 41.33 years.
6. Session Duration: The average session duration is slightly longer for users on the Light Theme (930.83 seconds) than for those on the Dark Theme (919.48 seconds).

From these insights, it appears that the Light Theme slightly outperforms the Dark Theme in terms of Conversion Rate, Bounce Rate, Scroll Depth, and Session Duration, while the Dark Theme leads in Click Through Rate. However, the differences are relatively minor across all metrics. Getting Started with Hypothesis Testing We'll use a significance level (alpha) of 0.05 for our hypothesis testing. It means we'll consider a result statistically significant if the p-value from our test is less than 0.05.

Let's start with hypothesis testing based on the Conversion Rate between the Light Theme and Dark Theme. Our hypotheses are as follows:

Null Hypothesis (H0): There is no difference in Conversion Rates between the Light Theme and Dark Theme. Alternative Hypothesis (Ha): There is a difference in Conversion Rates between the Light Theme and Dark Theme. We'll use a two-sample t-test to compare the means of the two independent samples. Let's proceed with the test:

```
[20]: # extracting conversion rates for both themes
conversion_rates_light = df[df['Theme'] == 'Light Theme']['Conversion Rate']
conversion_rates_dark = df[df['Theme'] == 'Dark Theme']['Conversion Rate']

[22]: # performing a two-sample t-test
t_stat, p_value = ttest_ind(conversion_rates_light, conversion_rates_dark,
                             equal_var=False)

t_stat, p_value
```

```
[22]: (0.4748494462782632, 0.6349982678451778)
```

The result of the two-sample t-test gives a p-value of approximately 0.635. Since this p-value is much greater than our significance level of 0.05, we do not have enough evidence to reject the null hypothesis. Therefore, we conclude that there is no statistically significant difference in Conversion Rates between the Light Theme and Dark Theme based on the data provided. Now, let's conduct hypothesis testing based on the Click Through Rate (CTR) to see if there's a statistically significant difference between the Light Theme and Dark Theme regarding how often users click through. Our hypotheses remain structured similarly:

Null Hypothesis (H0): There is no difference in Click Through Rates between the Light Theme and Dark Theme. Alternative Hypothesis (Ha): There is a difference in Click Rates between the Light Theme and Dark Theme.

```
[23]: # extracting click through rates for both themes
ctr_light = df[df['Theme'] == 'Light Theme']['Click Through Rate']
ctr_dark = df[df['Theme'] == 'Dark Theme']['Click Through Rate']

# performing a two-sample t-test
t_stat_ctr, p_value_ctr = ttest_ind(ctr_light, ctr_dark, equal_var=False)

t_stat_ctr, p_value_ctr
```

[23]: (-1.9781708664172253, 0.04818435371010704)

The two-sample t-test for the Click Through Rate (CTR) between the Light Theme and Dark Theme yields a p-value of approximately 0.048. This p-value is slightly below our significance level of 0.05, indicating that there is a statistically significant difference in Click Through Rates between the Light Theme and Dark Theme, with the Dark Theme likely having a higher CTR given the direction of the test statistic. Now, let's perform Hypothesis Testing based on two other metrics: bounce rate and scroll depth, which are important metrics for analyzing the performance of a theme or a design on a website. I'll first perform these statistical tests and then create a table to show the report of all the tests we have done

```
[24]: # extracting bounce rates for both themes
bounce_rates_light = df[df['Theme'] == 'Light Theme']['Bounce Rate']
bounce_rates_dark = df[df['Theme'] == 'Dark Theme']['Bounce Rate']

# performing a two-sample t-test for bounce rate
t_stat_bounce, p_value_bounce = ttest_ind(bounce_rates_light,
↪ bounce_rates_dark, equal_var=False)

# extracting scroll depths for both themes
scroll_depth_light = df[df['Theme'] == 'Light Theme']['Scroll_Depth']
scroll_depth_dark = df[df['Theme'] == 'Dark Theme']['Scroll_Depth']

# performing a two-sample t-test for scroll depth
t_stat_scroll, p_value_scroll = ttest_ind(scroll_depth_light,
↪ scroll_depth_dark, equal_var=False)

# creating a table for comparison
comparison_table = pd.DataFrame({
    'Metric': ['Click Through Rate', 'Conversion Rate', 'Bounce Rate', 'Scroll_
↪ Depth'],
    'T-Statistic': [t_stat_ctr, t_stat, t_stat_bounce, t_stat_scroll],
    'P-Value': [p_value_ctr, p_value, p_value_bounce, p_value_scroll]
})

comparison_table
```

[24]:	Metric	T-Statistic	P-Value
0	Click Through Rate	-1.978171	0.048184
1	Conversion Rate	0.474849	0.634998
2	Bounce Rate	-1.201888	0.229692
3	Scroll Depth	0.756228	0.449692

Click Through Rate: The test reveals a statistically significant difference, with the Dark Theme likely performing better (P-Value = 0.048). Conversion Rate: No statistically significant difference was found (P-Value = 0.635). Bounce Rate: There's no statistically significant difference in Bounce Rates between the themes (P-Value = 0.230). Scroll Depth: Similarly, no statistically significant difference is observed in Scroll Depths (P-Value = 0.450). In summary, while the two themes perform similarly across most metrics, the Dark Theme has a slight edge in terms of engaging users to click through. For other key performance indicators like Conversion Rate, Bounce Rate, and Scroll Depth, the choice between a Light Theme and a Dark Theme does not significantly affect user behaviour according to the data provided.

Summary

So, Hypothesis Testing is a statistical method used to make inferences or decisions about a population based on sample data. It starts with a null hypothesis (H_0), which represents a default stance or no effect, and an alternative hypothesis (H_1 or H_a), which represents what we aim to prove or expect to find. The process involves using sample data to determine whether to reject the null hypothesis in favor of the alternative hypothesis, based on the likelihood of observing the sample data under the null hypothesis.

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]: